



机器学习

苏州大学计算机科学与技术学院

自然语言处理实验室

主讲：周夏冰

邮箱：zhouxiabing@suda.edu.cn



01


BP神经网络



神经网络



- ◆模拟人脑神经系统的结构和功能，运用大量简单处理单元经广泛连接而组成的人工网络系统。
- ◆神经网络方法是一种**隐式**的知识表示方法
- ◆最早的神经网络的思想起源于1943年的MCP人工神经元模型

M-P模型



神经网络



- 第一次打破非线性诅咒的当属现代DL大牛Hinton，其在1986年发明了适用于多层感知器（MLP）的BP算法，并采用Sigmoid进行非线性映射，有效解决了非线性分类和学习的问题





BP神经网络


- 连接主义的神经网络有着多种多样的网络结构以及学习方法，虽然早期模型强调模型的生物可解释性（Biological Plausibility），但后期更关注于对某种特定认知能力的模拟，比如物体识别、语言理解等. 尤其在引入**误差反向传播**来改进其学习能力之后，神经网络也越来越多地应用在各种机器学习任务上.
- BP（back-propagation）神经网络就是**多层前向网络**，利用**反向误差传播**来进行学习



BP学习算法

• 前向传播

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$





$$1 \times 1 + (-1) \times (-1) + 1 = 3$$

$$\frac{1}{1 + e^{-3}} = 0.95$$

$$1 \times 1 + (-1) \times 2 - 1 = -2 \quad 1 \times (-1) + (-1) \times 1 - 0 = -2$$

$$\frac{1}{1 + e^2} = 0.12$$


$$\frac{1}{1 + e^2} = 0.12$$



BP学习算法

• 前向传播

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



$$0.12 \times 2 + 0.12 \times 1 + 0.95 \times 4 - 0 = 4.16$$




$$\frac{1}{1 + e^{-4.16}} = 0.98$$

$$0.12 \times 3 + 0.12 \times (-2) + 0.95 \times (-1) + 1 = 0.39$$

$$\frac{1}{1 + e^{-0.39}} = 0.6$$



BP学习算法







BP学习算法


■ 误差反向传播

■ $\widehat{w}_i \leftarrow w_i - \eta \frac{\Delta E}{\Delta w_i}$





BP学习算法

■ 向前传播



■ 对于训练样例: (x_k, y_k)


➤ 均方误差: $E_k = \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2$



BP学习算法

反向传播

➤ 梯度下降



输出层损失函数

➤ $L = \frac{1}{2} (\hat{y}_k^j - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} \alpha_h - \theta_2) - y_k^j)^2$

➤ $\frac{\partial L}{\partial w_{hj}} = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial w_{hj}} = (\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j)) \cdot \alpha_h$

$$f = \frac{1}{1 + e^{-z}}$$


$$f' = -\frac{1}{(1 + e^{-z})^2} \cdot (-e^{-z})$$

$$= \frac{e^{-z}}{(1 + e^{-z})^2} = \frac{1 + e^{-z} - 1}{(1 + e^{-z})^2}$$

$$= \frac{1}{1 + e^{-z}} \cdot \frac{1 + e^{-z} - 1}{1 + e^{-z}}$$

$$= f \cdot (1 - f)$$

BP学习算法



$$\beta_j = \sum_{h=1}^q w_{hj} \alpha_h - \theta_2$$


$$\gamma_h = \sum_{i=1}^d v_{ih} x^i - \theta_1$$

$$\frac{\partial L}{\partial w_{hj}} = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial w_{hj}} = \underbrace{(\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j))}_{\delta_j^L} \cdot \alpha_h \quad \delta_j^L = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} = \frac{\partial L}{\partial \beta_j}$$

$$\begin{aligned} L &= \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\beta_j) - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} \alpha_h - \theta_2) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\gamma_h))) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\sum_{i=1}^d v_{ih} x^i - \theta_1)) - \theta_2) - y_k^j)^2 \end{aligned}$$

$$\frac{\partial L}{\partial v_{ih}} = \sum_{j=1}^l (\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j)) \cdot w_{hj} \cdot f(\gamma_h)(1 - f(\gamma_h)) \cdot x^i$$

BP学习算法



$$\beta_j = \sum_{h=1}^q w_{hj} \alpha_h - \theta_2$$


$$\gamma_h = \sum_{i=1}^d v_{ih} x^i - \theta_1$$

$$\frac{\partial L}{\partial w_{hj}} = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial w_{hj}} = \underbrace{(\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j))}_{\delta_j^L} \cdot \alpha_h \quad \delta_j^L = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} = \frac{\partial L}{\partial \beta_j}$$

$$\begin{aligned} L &= \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\beta_j) - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} \alpha_h - \theta_2) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\gamma_h))) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\sum_{i=1}^d v_{ih} x^i - \theta_1)) - \theta_2) - y_k^j)^2 \end{aligned}$$


$$\frac{\partial L}{\partial v_{ih}} = \underbrace{\sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot f(\gamma_h)(1 - f(\gamma_h))}_{\delta_h^{L-1}} \cdot x^i \quad \delta_h^{L-1} = \sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot \frac{\partial \alpha_h}{\partial \gamma_h}$$

BP学习算法



$$\delta_h^{L-1} = \sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot \frac{\partial \alpha_h}{\partial \gamma_h}$$

BP学习算法



$$\beta_j = \sum_{h=1}^q w_{hj} \alpha_h - \theta_2$$

$$\gamma_h = \sum_{i=1}^d v_{ih} x^i - \theta_1$$

$$\frac{\partial L}{\partial w_{hj}} = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial w_{hj}} = \underbrace{(\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j))}_{\delta_j^L} \cdot \alpha_h$$

$$\delta_j^L = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} = \frac{\partial L}{\partial \beta_j}$$

$$\begin{aligned} L &= \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\beta_j) - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} \alpha_h - \theta_2) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\gamma_h))) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\sum_{i=1}^d v_{ih} x^i - \theta_1)) - \theta_2) - y_k^j)^2 \end{aligned}$$


$$\frac{\partial L}{\partial v_{ih}} = \underbrace{\sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot f(\gamma_h)(1 - f(\gamma_h))}_{\delta_h^{L-1}} \cdot x^i$$

$$\delta_h^{L-1} = \sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot \frac{\partial \alpha_h}{\partial \gamma_h} = \sum_{j=1}^l \frac{\partial L}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial \alpha_h} \cdot \frac{\partial \alpha_h}{\partial \gamma_h} = \frac{\partial L}{\partial \gamma_h}$$

$$= \frac{\partial L}{\partial \alpha_h} \frac{\partial \alpha_h}{\partial \gamma_h}$$

BP学习算法

前提 $L = \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2$



$$\beta_j = \sum_{h=1}^q w_{hj} \alpha_h - \theta_2$$

$$\gamma_h = \sum_{i=1}^d v_{ih} x^i - \theta_1$$

$$\frac{\partial L}{\partial w_{hj}} = \frac{\partial L}{\partial \hat{y}_k^j} \cdot \frac{\partial \hat{y}_k^j}{\partial \beta_j} \cdot \frac{\partial \beta_j}{\partial w_{hj}} = \underbrace{(\hat{y}_k^j - y_k^j) \cdot f(\beta_j)(1 - f(\beta_j))}_{\delta_j^L} \cdot \alpha_h = \delta_j^L \cdot \alpha_h$$


$$\begin{aligned} L &= \frac{1}{2} \sum_{j=1}^l (\hat{y}_k^j - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\beta_j) - y_k^j)^2 = \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} \alpha_h - \theta_2) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\gamma_h)) - y_k^j)^2 \\ &= \frac{1}{2} \sum_{j=1}^l (f(\sum_{h=1}^q w_{hj} (f(\sum_{i=1}^d v_{ih} x^i - \theta_1) - \theta_2) - y_k^j)^2 \end{aligned}$$

$$\frac{\partial L}{\partial v_{ih}} = \delta_h^{L-1} \cdot x^i \quad \delta_h^{L-1} = \sum_{j=1}^l \delta_j^L \cdot w_{hj} \cdot f(\gamma_h)(1 - f(\gamma_h))$$



BP学习算法

- 共有L层
- 第 l 层神经元共有 n_l
- 第 $l-1$ 层到 l 层的权重为: w_{ij}^l
- 第 l 层的输入为: h_i^l



$$\frac{\partial E}{\partial w_{ij}^L} = \delta_j^L \cdot h_i^{L-1}$$

$$\delta_j^L = (\hat{y}_k^j - y_k^j)(f_j^L)'$$




BP学习算法



$$\frac{\partial E}{\partial w_{ij}^l} = \delta_j^l \cdot h_i^{l-1}$$

$$\delta_j^l = \sum_{s=1}^{n_{l+1}} \delta_s^{l+1} \cdot w_{js}^{l+1} \cdot (f_j^l)'$$



BP学习算法

所有参数初始化 w, θ

for iter=1 to max (迭代次数)

for k=1 to M: (所有样本都要计算)

前向传播过程, 计算出所有的输出值 $\hat{y}_k^1, \dots, \hat{y}_k^{n_L}$, 并且保留所有隐层输入值 h

//计算每一层的 Δw $\frac{\partial E_k}{\partial \theta_j^L} = \delta_j^L$

$$\frac{\partial E_k}{\partial w_{ij}^L} = \delta_j^L \cdot h_i^{L-1}; \delta_j^L = (\hat{y}_k^j - y_k^j)(f_j^L)'; f_j^L = \text{sigmoid}(\sum_{d=1}^{n_{L-1}} h_d^{L-1} w_{dj}^L - \theta_j^L)$$

for l=l-1 to 2 $\frac{\partial E}{\partial \theta_j^l} = \delta_j^l$

$$\frac{\partial E}{\partial w_{ij}^l} = \delta_j^l \cdot h_i^{l-1}; \delta_j^l = \sum_{s=1}^{n_{l+1}} \delta_s^{l+1} \cdot w_{js}^{l+1} \cdot (f_j^l)'; f_j^l = \text{sigmoid}(\sum_{d=1}^{n_{l-1}} h_d^{l-1} w_{dj}^l - \theta_j^l)$$

//更新每一层的权值

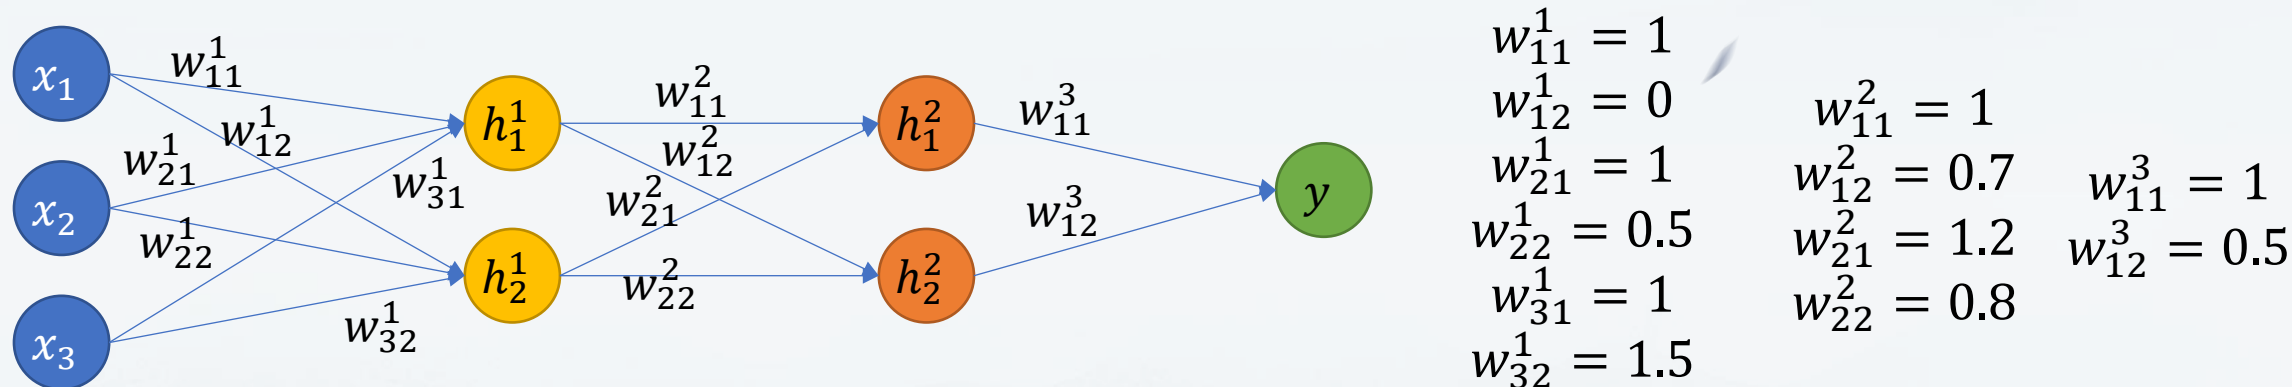
for l=2 to L

$$w_{ij}^l = w_{ij}^l - \eta \sum_{k=1}^M \frac{\partial E_k}{\partial w_{ij}^l}$$

课堂练习

- 如下的BP神经网络，学习步长 $\eta = 1$ ，各点的阈值 $\theta = 0$ ，输入层到隐层的激活函数为

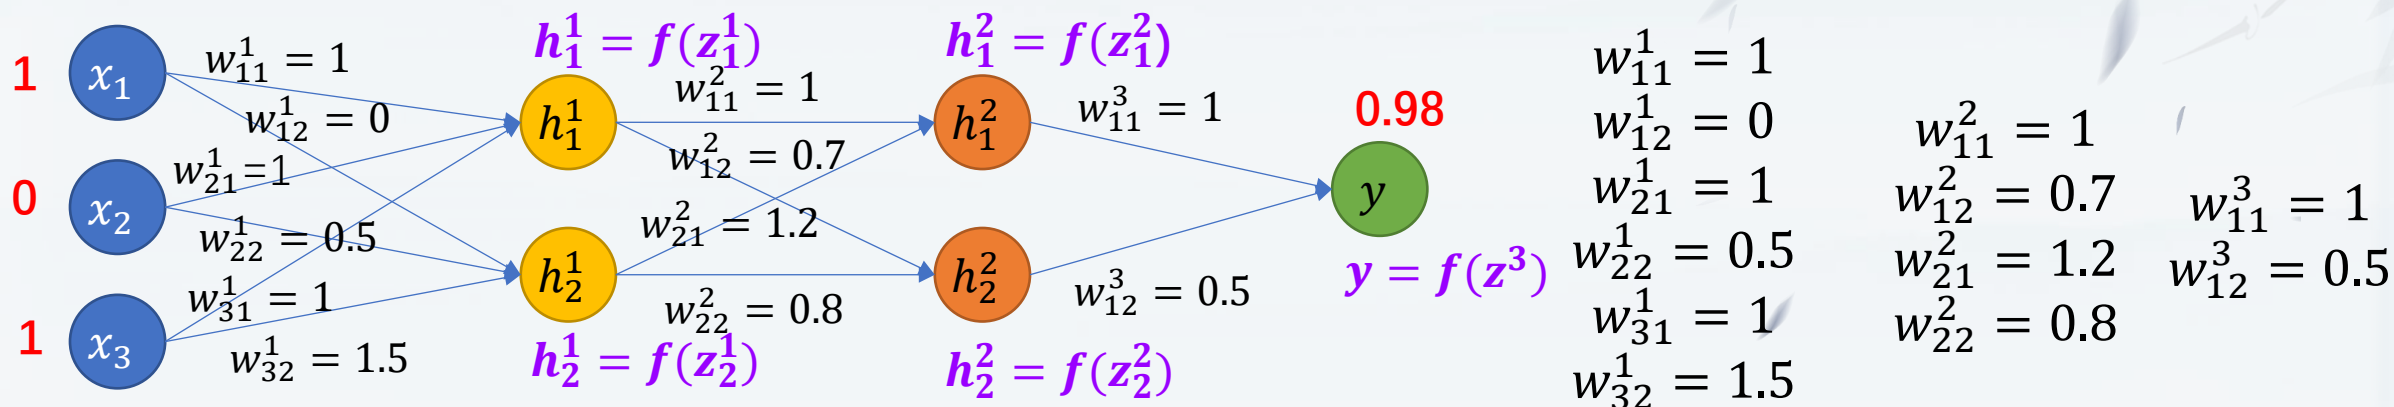
$f(x) = \max(1, x)$ ，隐层到隐层以及隐层到输出层的激活函数为 $f(x) = \frac{1}{1+e^{-x}}$



设输入样本 $x_1 = 1, x_2 = 0, x_3 = 1$ ，输出节点的期望输出 $y = 0.98$

利用预测误差 $E = \frac{1}{2}(\hat{y} - y)^2$ 对连接权进行调整（只调整一轮，阈值更新不计算）

课堂练习



前向传播

- $z_1^1 = 1 * 1 + 0 * 1 + 1 = 2, \quad z_2^1 = 1 * 0 + 0 * 0.5 + 1 * 1.5 = 1.5$

- $h_1^1 = \max(1, 2) = 2, \quad h_2^1 = \max(1, 1.5) = 1.5$

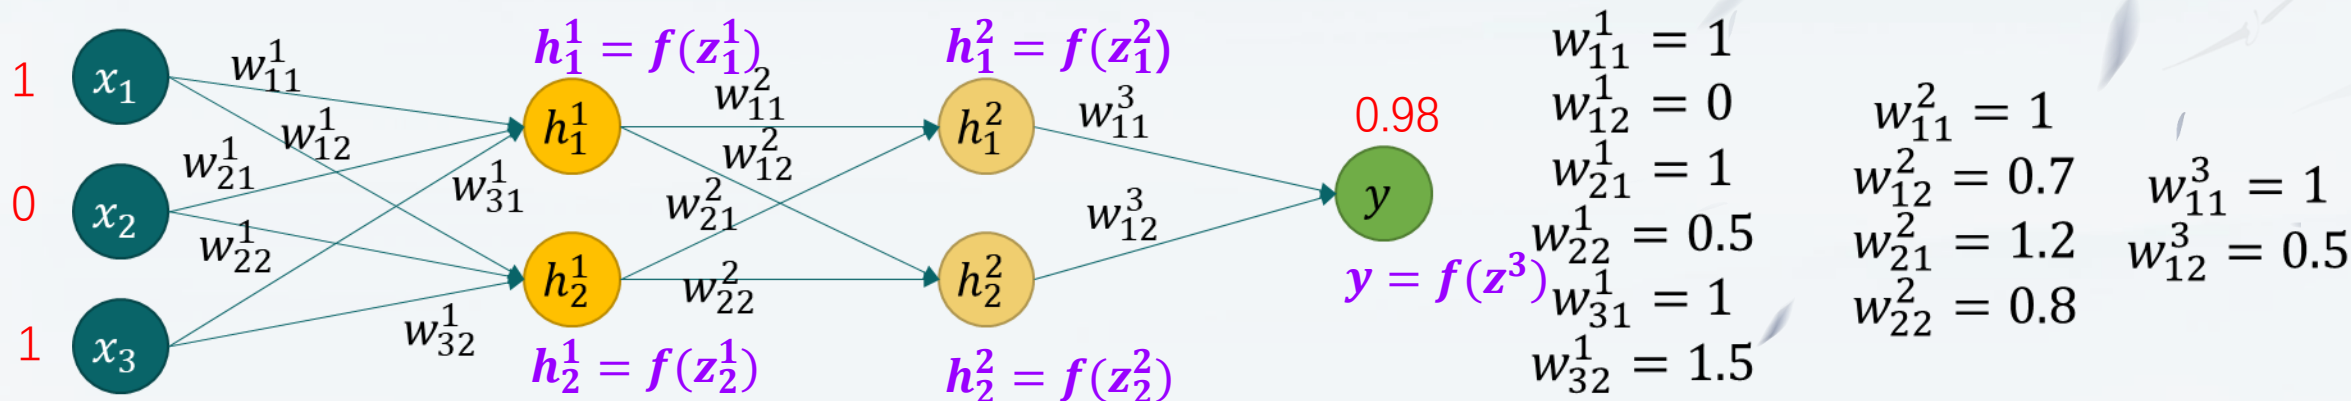
- $z_1^2 = 2 * 1 + 1.5 * 1.2 = 3.8, \quad z_2^2 = 2 * 0.7 + 1.5 * 0.8 = 2.6$

- $h_1^1 = \frac{1}{1+e^{-3.8}} = 0.9781, \quad h_2^1 = \frac{1}{1+e^{-2.6}} = 0.9309$

$$z^3 = 0.9781 * 1 + 0.9309 * 0.5 = 1.4436$$

$$\hat{y} = \frac{1}{1+e^{-1.4436}} = 0.809$$

课堂练习



反向传播

$$f = \frac{1}{1 + e^{-z}}$$

$$\delta_j^L = \frac{\partial L}{\partial z_j^L} = \frac{\partial L}{\partial y_j} \frac{\partial y_j}{\partial z_j^L}$$

$$\delta_1^3 = (\hat{y} - y) * f'(z^3) = (0.809 - 0.98) * 0.809 * (1 - 0.809) = -0.0264$$


$$\delta_1^2 = \delta_1^3 * w_{11}^3 * f'(z_1^2) = -0.0264 * 1 * 0.9781 * (1 - 0.9781) = -0.0006$$

$$\delta_2^2 = \delta_1^3 * w_{12}^3 * f'(z_2^2) = -0.0264 * 0.5 * 0.9309 * (1 - 0.9309) = -0.0008$$

$$\delta_1^1 = (\delta_1^2 * w_{11}^2 + \delta_2^2 * w_{12}^2) * f'(z_1^1) = (-0.0006 * 1 - 0.0008 * 0.7) * 1 = -0.0012$$

$$\delta_2^1 = (\delta_1^2 * w_{21}^2 + \delta_2^2 * w_{22}^2) * f'(z_2^1) = (-0.0006 * 1.2 - 0.0008 * 0.8) * 1 = -0.0014$$

课堂练习



$w_{11}^1 = 1$	$w_{11}^2 = 1$	$w_{11}^3 = 1$
$w_{12}^1 = 0$	$w_{12}^2 = 0.7$	$w_{12}^3 = 0.5$
$w_{21}^1 = 1$	$w_{21}^2 = 1.2$	
$w_{22}^1 = 0.5$	$w_{22}^2 = 0.8$	
$w_{31}^1 = 1$		
$w_{32}^1 = 1.5$		

反向传播

$$\Delta w_{11}^3 = \delta_1^3 * h_1^2 = -0.0264 * 0.9781 = -0.0258$$

$$\Delta w_{12}^3 = \delta_1^3 * h_2^2 = -0.0264 * 0.9309 = -0.0246$$

$$\Delta w_{11}^2 = \delta_1^2 * h_1^1 = -0.0006 * 2 = -0.0012$$

$$\Delta w_{12}^2 = \delta_2^2 * h_1^1 = -0.0008 * 2 = -0.0016$$

$$\Delta w_{21}^2 = \delta_1^2 * h_1^1 = -0.0006 * 1.5 = -0.0009$$

$$\Delta w_{22}^2 = \delta_2^2 * h_1^1 = -0.0008 * 1.5 = -0.0012$$

$$\Delta w_{11}^1 = \delta_1^1 * x_1 = -0.0012 * 1 = -0.0012$$

$$\Delta w_{12}^1 = \delta_2^1 * x_1 = -0.0014 * 1 = -0.0014$$

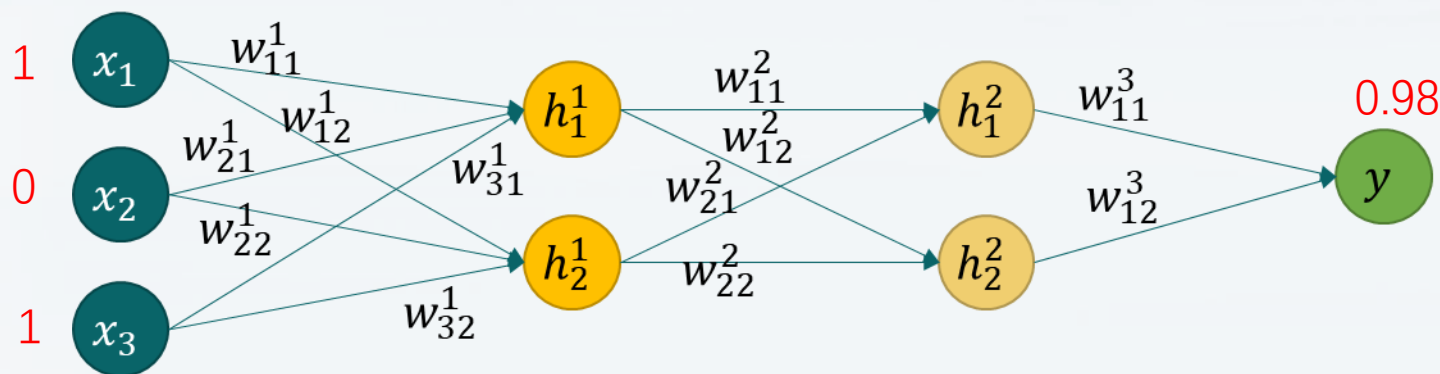
$$\Delta w_{21}^1 = \delta_1^1 * x_2 = -0.0012 * 0 = -0.0012$$

$$\Delta w_{22}^1 = \delta_2^1 * x_2 = -0.0014 * 0 = -0.0014$$

$$\Delta w_{31}^1 = \delta_1^1 * x_3 = -0.0012 * 1 = -0.0012$$

$$\Delta w_{32}^1 = \delta_2^1 * x_3 = -0.0014 * 1 = -0.0014$$

课堂练习



$$\begin{aligned}w_{11}^1 &= 1 \\w_{12}^1 &= 0 \\w_{21}^1 &= 1 \\w_{22}^1 &= 0.5 \\w_{31}^1 &= 1 \\w_{32}^1 &= 1.5\end{aligned}$$

$$\begin{aligned}w_{11}^2 &= 1 \\w_{12}^2 &= 0.7 \\w_{21}^2 &= 1.2 \\w_{22}^2 &= 0.8\end{aligned}$$

$$\begin{aligned}w_{11}^3 &= 1 \\w_{12}^3 &= 0.5\end{aligned}$$

反向传播

$$\begin{aligned}w_{11}^1 &= 1 + 1 * 0.0012 = 1.0012 \\w_{12}^1 &= 0 + 0.0014 = 0.0014 \\w_{21}^1 &= 1 + 0 = 1 \\w_{22}^1 &= 0.5 + 0 = 0.5 \\w_{31}^1 &= 1 + 0.0012 = 1.0012 \\w_{32}^1 &= 1.5 + 0.0014 = 1.5014\end{aligned}$$

$$\begin{aligned}w_{11}^2 &= 1 + 0.0012 = 1.0012 \\w_{12}^2 &= 0.7 + 0.0016 = 0.7016 \\w_{21}^2 &= 1.2 + 0.0009 = 1.2009 \\w_{22}^2 &= 0.8 + 0.0012 = 0.8012\end{aligned}$$

$$\begin{aligned}w_{11}^3 &= 1 + 0.0258 = 1.0258 \\w_{12}^3 &= 0.5 + 0.0246 = 0.5246\end{aligned}$$



神经网络



- 在1989年以后由于没有特别突出的方法被提出，且NN一直缺少相应的严格的数学理论支持，神经网络的热潮渐渐冷淡下去。冰点来自于1991年，BP算法被指出存在梯度消失问题

Sigmoid函数: $\sigma(x) = \frac{1}{1+\exp(-x)}$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$



饱和区的导数接近0



误差传递不断衰减

梯度消失

神经网络



- 在1989年以后由于没有特别突出的方法被提出，且NN一直缺少相应的严格的数学理论支持，神经网络的热潮渐渐冷淡下去。冰点来自于1991年，BP算法被指出存在梯度消失问题
- 2006年，Hinton提出了深层网络训练中梯度消失问题的解决方案
- 2011年，ReLU激活函数被提出，该激活函数能够有效的抑制梯度消失问题



激活函数




• ReLU函数

$$\bullet \text{ReLU}(x) = \begin{cases} x & x \geq 0 \\ 0 & x < 0 \end{cases}$$

- 1) 计算更加高效
- 2) 求导在 $x > 0$ 时为1，一定程度缓解梯度消失现象

如果出现不恰当更新，某个神经元不被激活，自身梯度永远为0：死亡ReLU问题



Optimization of New Model

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} L$$

➤ (Randomly) Pick initial values $\boldsymbol{\theta}^0$

➤ Compute gradient $\boldsymbol{g} = \nabla L(\boldsymbol{\theta}^0)$



$$\boldsymbol{\theta}^1 \leftarrow \boldsymbol{\theta}^0 - \eta \boldsymbol{g}$$

➤ Compute gradient $\boldsymbol{g} = \nabla L(\boldsymbol{\theta}^1)$



$$\boldsymbol{\theta}^2 \leftarrow \boldsymbol{\theta}^1 - \eta \boldsymbol{g}$$

➤ Compute gradient $\boldsymbol{g} = \nabla L(\boldsymbol{\theta}^2)$

$$\boldsymbol{\theta}^3 \leftarrow \boldsymbol{\theta}^2 - \eta \boldsymbol{g}$$



Optimization of New Model



critical point



The background of the slide is a soft-focus photograph of pink cherry blossoms. The flowers are in various stages of bloom, with delicate petals and visible yellow stamens. The lighting is bright and natural, creating a warm and pleasant atmosphere.

Batch



Optimization of New Model

$$\theta^* = \arg \min_{\theta} L$$

➤ (Randomly) Pick initial values θ^0

➤ Compute gradient $g = \nabla L^1(\theta^0)$


$$\text{update } \theta^1 \leftarrow \theta^0 - \eta g$$

➤ Compute gradient $g = \nabla L^2(\theta^1)$

$$\text{update } \theta^2 \leftarrow \theta^1 - \eta g$$

➤ Compute gradient $g = \nabla L^3(\theta^2)$

$$\text{update } \theta^3 \leftarrow \theta^2 - \eta g$$



1 **epoch** = see all the batches once


Batch



Consider 20 examples ($N=20$)


Batch size = N (Full batch)

Update after seeing all
the 20 examples



Batch size = 1


Update for each example
Update 20 times in an epoch






Batch



Full Batch



Small Batch




The background of the slide is a soft-focus photograph of pink cherry blossoms. The flowers are in various stages of bloom, with delicate petals and visible yellow stamens. The lighting is bright and natural, creating a gentle, spring-like atmosphere.


Momentum



Momentum



(Vanilla) Gradient Descent



初始 θ^0

梯度计算 g^0


$$\theta^1 = \theta^0 - \eta g^0$$

梯度计算 g^1

$$\theta^2 = \theta^1 - \eta g^1$$

...

Gradient Descent + Momentum



初始 θ^0

设 $m^0 = 0$

梯度计算 g^0

$$m^1 = \lambda m^0 - \eta g^0$$


$$\theta^1 = \theta^0 + m^1$$

梯度计算 g^1

$$m^2 = \lambda m^1 - \eta g^1$$

$$\theta^2 = \theta^1 + m^2$$

Gradient Descent + Momentum




The background of the slide is a soft-focus photograph of pink cherry blossoms. The flowers are in various stages of bloom, with delicate petals and visible yellow stamens. The lighting is bright and natural, creating a gentle, spring-like atmosphere.


Learning rate




Training stuck \neq Small Gradient





loss




norm of gradient






$$\eta = 10^{-2}$$



$$\eta = 10^{-7}$$



Learning rate



$$\theta_i^{t+1} \leftarrow \theta_i^t - \boxed{\eta} g_i^t$$

$$g_i^t = \frac{\partial L}{\partial \theta_i} \bigg|_{\theta = \theta^t}$$

$$\theta_i^{t+1} \leftarrow \theta_i^t - \boxed{\frac{\eta}{\sigma_i^t}} g_i^t$$






Learning rate

- $\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$

- $\sigma_i^t = \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g_i^t)^2}$

- Adagrad



Learning rate

- $\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$


- $\sigma_i^t = \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g_i^t)^2}$

- **Adagrad**

- $\sigma_i^t = \sqrt{\alpha (\sigma_i^{t-1})^2 + (1 - \alpha) (g_i^t)^2}$


- **RMSProp**

Adam: 结合了动量梯度下降和RMSProp两种算法的优点




Learning rate



$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$



Learning Rate Decay



Warm Up



Summary of Optimization



(Vanilla) Gradient Descent

$$\theta_i^{t+1} \leftarrow \theta_i^t - \eta g_i^t$$

Various Improvements

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta^t}{\sigma_i^t} m_i^t$$

Learning rate scheduling

Momentum: weighted sum of the previous gradients

root mean square of the gradients

确定方向

确定大小

